

First tutorial session

Vincent Dumoulin

January 15, 2015

- 1 Solution to the numpy + MNIST + MLP assignment
- 2 Git primer
- 3 Theano primer
- 4 Porting numpy + MNIST + MLP to theano + MNIST + MLP

Solution to numpy + MNIST + MLP

Setup

$$\mathbf{x} = [x_1, \dots, x_A]^T,$$

$$\mathbf{t} = [t_1, \dots, t_C]^T,$$

$$\mathbf{W} = \begin{pmatrix} W_{1,1} & \cdots & W_{1,A} \\ \cdots & \cdots & \cdots \\ W_{H,1} & \cdots & W_{H,A} \end{pmatrix}, \quad \mathbf{b} = [b_1, \dots, b_H]^T,$$

$$\mathbf{V} = \begin{pmatrix} V_{1,1} & \cdots & V_{1,H} \\ \cdots & \cdots & \cdots \\ V_{C,1} & \cdots & V_{C,H} \end{pmatrix}, \quad \mathbf{c} = [c_1, \dots, c_C]^T,$$

$$\mathbf{h} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b}),$$

$$\mathbf{y} = \text{softmax}(\mathbf{V}\mathbf{h} + \mathbf{c})$$

$$\mathcal{L} = -\mathbf{t}^T \log \mathbf{y}$$

Derivatives

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial \mathbf{V}} &= (\mathbf{y} - \mathbf{t})\mathbf{h}^T, \\ \frac{\partial \mathcal{L}}{\partial \mathbf{c}} &= \mathbf{y} - \mathbf{t} \\ \frac{\partial \mathcal{L}}{\partial \mathbf{W}} &= [\mathbf{V}^T(\mathbf{y} - \mathbf{t}) \odot \mathbf{h} \odot (\mathbf{1} - \mathbf{h})]\mathbf{x}^T, \\ \frac{\partial \mathcal{L}}{\partial \mathbf{b}} &= \mathbf{V}^T(\mathbf{y} - \mathbf{t}) \odot \mathbf{h} \odot (\mathbf{1} - \mathbf{h})\end{aligned}$$

A complete derivation can be found here

<https://raw.githubusercontent.com/vdumoulin/ift6266h15/master/assignments/01/solution.pdf>

Solution to numpy + MNIST + MLP

The full python implementation can be found here

```
https://raw.githubusercontent.com/vdumoulin/  
ift6266h15/master/assignments/01/solution.py
```

Why version control?

- A nice and clean alternative to maintaining multiple versions of the same file using some sort of custom naming scheme (e.g. `myfile_9.py`)
- A way to end the fear of saving and quitting
- Keep a trace how your files change throughout development
- Revert back to older versions
- Manage multiple versions (*branches*) of your code at the same time

Going further

<http://git-scm.com/book/en/v2/Getting-Started-About-Version-Control>

What is git?

- Distributed version control system
 - No checking out single files: local version fully mirrors the repository
 - No central authority on what is the *true* codebase
- Takes *snapshots* of the state of a repository at a given time
- Intelligent about not duplicating information from one snapshot to another

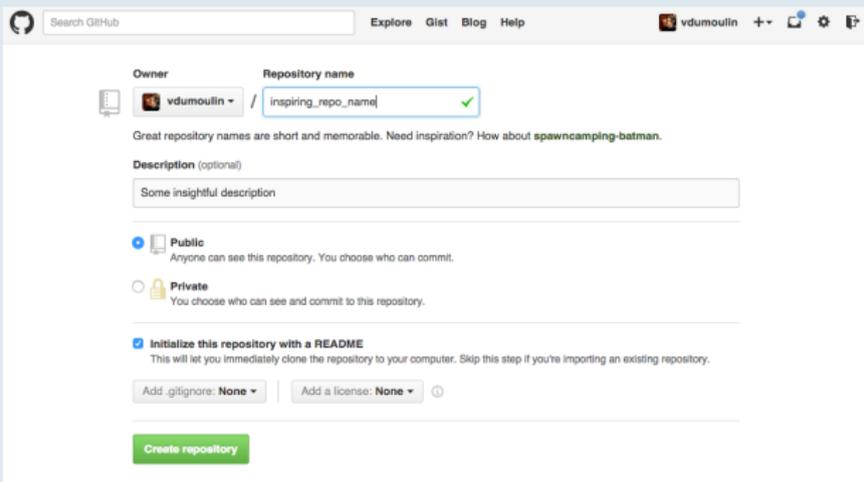
Going further

<http://git-scm.com/book/en/v2/Getting-Started-Git-Basics>

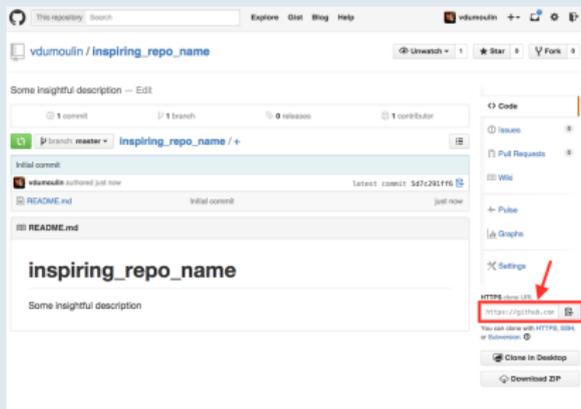
What is Github?

- A place to host your git repositories
- Makes it easy to
 - Share code with others
 - Keep track of other people's code
 - Modify other people's code (*forking*)
 - Collaborate with other people on common code
- Technically no different from your own machine:
 - Both can pull and push changes
 - Both host a fully functional version of your repository

Creating a repository on Github



Identifying the repository URL



The screenshot shows a GitHub repository page for 'vdmoulin / inspiring_repo_name'. The repository is currently on the 'master' branch. The page includes a search bar, repository statistics (1 commit, 1 branch, 0 releases, 1 contributor), and a list of files including 'README.md'. On the right-hand side, under the 'Code' section, the 'HTTPS clone URL' is displayed as 'https://github.com/vdmoulin/inspiring_repo_name.git'. This URL is highlighted with a red box, and a red arrow points to it from the 'Settings' link above.

Cloning a git repository

```
> git clone <REPO URL>
```

Putting a file under version control

- Create a dummy file
- Check the status of the repository:

```
> git status
```

- Add the file to version control:

```
> git add my_dummy_file.py
```

- Commit the newly-added file:

```
> git commit -m "Add_new_dummy_file_to_repository"
```

Commit changes to a file

- Change something in your file
- Stage the changes:

```
> git add my_dummy_file.py
```

- Commit the changes:

```
> git commit -m "Add_stuff_to_dummy_file"
```

Pull changes on Github

- Run

```
> git pull origin master
```

Push changes on Github

- Pull the latest changes from your Github repo:

```
> git pull origin master
```

- Push your changes to Github:

```
> git push origin master
```

Going further with Git

<http://git-scm.com/book/en/v2>

What is Theano?

- From Theano's online documentation:

Theano is a Python library that allows you to define, optimize, and evaluate mathematical expressions involving multi-dimensional arrays efficiently.

- Does *symbolic* computation **and** differentiation (*i.e.* the end result of differentiation is itself a symbolic expression)
- Very similar to numpy with respect to its interface
- Allows doing numerical computation in a high-level language (Python) while still retaining the speed of low-level languages (like C)
- Allows the generation of efficient CPU and GPU code transparently

Typical Theano workflow

- 1 Instantiate symbolic variables
- 2 Build a computation graph out of those variables
- 3 Compile a function with the symbolic variables as input and the output of the computation graph as output
- 4 Call the compiled function with numerical inputs

Theano vs. numpy

- Theano interface is *very* similar to numpy interface
- numpy arrays are automatically converted to constant symbolic variables when used inside a computation graph
- You can manipulate Theano symbolic variables in the same way you'd manipulate numpy arrays

Going further: Theano's basic interface

<http://deeplearning.net/software/theano/library/tensor/basic.html>

Types of symbolic variables

TensorVariable Its value is unspecified at graph creation and can change from one call of the compiled function to another (e.g. x and y in $y = 3x - 2$). **Not persistent across function calls**

TensorConstant Its value is specified at graph creation and does **not** change from one call of the compiled function to another (e.g. 3 and -2 in $y = 3x - 2$)

TensorSharedVariable Its value is specified at graph creation but is bound to change from one call of the compiled function to another (e.g. a and b in $y = ax + b$ in a regression setting where some x and y pairs have been observed). **Persistent across function calls**

Examples

Listing 1: Simple algebra

```
import theano
import theano.tensor as T

# 1. Instantiate symbolic variables
x = T.vector(name='x')
y = T.vector(name='y')

# 2. Build a computation graph
z = x + y

# 3. Compile a callable function
f = theano.function(inputs=[x, y], outputs=z)

# 4. Call the function using numerical inputs
print f([1, 2], [3, 4])
```

Examples

Listing 2: Gradient computation

```
import theano
import theano.tensor as T

# 1. Instantiate symbolic variables
x = T.vector(name='x')

# 2. Build a computation graph
z = (x ** 2).sum()
d_z_d_x = T.grad(z, x)

# 3. Compile a callable function
f = theano.function(inputs=[x], outputs=d_z_d_x)

# 4. Call the function using numerical inputs
print f([1, 2])
```

Examples

Listing 3: Linear regression

```
import theano
import theano.tensor as T

x = T.scalar(name='x'); t = T.scalar(name='t')
a = theano.shared(-1.0, name='a')
b = theano.shared(0.0, name='b')

y = a * x + b
mse = (y - t) ** 2
grad_a, grad_b = T.grad(mse, [a, b])

f = theano.function(inputs=[x, t], outputs=mse,
                    updates={a: a - 0.01 * grad_a,
                              b: b - 0.01 * grad_b})

print [f(1, 5) for i in xrange(10)]
```

Going further: online Theano tutorial

```
http://deeplearning.net/software/theano/  
tutorial/index.html#tutorial
```